

Workshop: The V-Score Validator Capstone

Session 7 · Team Capstone · 30–35 min · 4 teams · 2–3 per team

Today is **execution, not lecture**. Don't build a calculator —
build a **team of scorer agents** that reason out each score, and orchestrate
them.

Setup was handled back in Session 2 — no separate setup check today. Breakout rooms, straight to the brief.

The v1 → v2 Shift

Same idea, scored. But the scoring is now done by **agents**, not arithmetic.

v1 · CALCULATOR

- Human **types 1–10** per criterion
- Code **multiplies** the numbers
- Your job: wire up arithmetic
- Judged on: does the math work

v2 · AGENT TEAM

- Human writes the **idea in plain English**
- A skill/agent per criterion **reasons out** each score
- Your job: **decompose** → **orchestrate** → **verify**
- Judged on: is the orchestration well-designed

| *The math is trivial and unchanged. The challenge is the agent system around it.*

What You Build

Two component types. Given one plain-language idea → two scores, a verdict, and a reason per criterion.



A Scorer per criterion

A SKILL.md or sub-agent. Takes `{ idea, rubric }` → returns `{ score 1-10, reason }`.


| Session 3: the *description* is what makes it fire.



An Orchestrator

Reads the idea, fans out to scorers, applies weights, sums, and maps to a verdict.

| Session 5: orchestrator + specialists.



Scoring Rubric: PoC + Market Viability

What your **scorers** evaluate — each returns a 1–10 with a reason, for both axes.

POC VIABILITY

Technical Novelty — new / unproven?	× 3
Defined Scope — narrowly definable?	× 4
Resource Accessibility — tools ready?	× 2
Measurable Outcome — success visible?	× 1

$$\text{PoC} = (\text{Novelty} \times 3) + (\text{Scope} \times 4) + (\text{Resources} \times 2) + (\text{Outcome} \times 1)$$

MARKET VIABILITY

Pain Severity — how acute?	× 4
Willingness to Pay — will they pay?	× 3
Market Size — big enough?	× 2
Differentiation — defensible?	× 1

$$\text{Market} = (\text{Pain} \times 4) + (\text{Pay} \times 3) + (\text{Size} \times 2) + (\text{Moat} \times 1)$$

Both max at 100. Never score a number without a reason behind it.

The Recommendation Matrix

Score ≥ 65 = **High**. The orchestrator maps the two axes to one verdict.

High PoC · High Market

 **Go / Full Speed Ahead**

Excellent on both axes.

Low PoC · High Market

 **De-risk First**

Strong demand, hard build. Spike the tech.

High PoC · Low Market

 **Validate Demand**

Buildable — but prove someone wants it.

Low PoC · Low Market

 **Reframe or Shelf**

High risk on both axes.

Print the verdict in **plain language** — plus the per-criterion reasons.

Pick Your Orchestration Pattern

Applying what you learned in Session 5 — choose deliberately, be ready to defend it in the demo.

SIMPLEST

Single agent + skills

One agent loads a scorer skill per criterion. Progressive disclosure. Great for 30 min.

MOST AGENTIC

Orchestrator + sub-agents

A coordinator spawns one scorer sub-agent per criterion, each blind to the others.

FASTEST

Parallel scorers

Fan all 8 scorers out at once, then aggregate. Show the parallelism live.

Deliverables — Bring 4 Things

Present in order in your 3-min demo. This is the scored part.

- 1 **The Spec** — brief + agent breakdown + measurable Definition of Done, written before building.
- 2 **The Agent/Skill definitions** — your `SKILL.md` (s) / sub-agent prompts. Point at the description that makes each fire.
- 3 **The Orchestration** — how the coordinator fans out, aggregates, weights, and decides.
- 4 **The Proof** — a live run on one new idea the judges hand you: per-criterion reasons, both scores, verdict.

Judging Rubric — 25 pts + 5

Weight is on **agentic design and steering**, not the calculator.

Spec quality — narrow scope, agent breakdown, measurable DoD	5
Agentic decomposition — a real skill/agent per criterion, sharp descriptions	5
Orchestration — coordinator fans out, aggregates, weights, decides; flow is legible	5
Verification & correctness — math + matrix correct; verified on a fresh idea	5
Demo clarity — the orchestration story lands in 3 min, reasons visible	5
Advanced bonus — memory bank and/or self-learning (demoed live)	+5

Scale per dimension: **0** absent · **2** weak · **4** solid · **5** exemplary → team total max **30**.

Advanced Tier — Memory & Self-Learning

Optional · +5 · after the core system works. Pick either or both.

A · Memory Bank

The orchestrator **remembers every idea** scored.

- Append each run to `scores.md` / `.json`
- On a new idea, recall the **most-similar past outcome**

| *Session 4 — demo the file growing.*

B · Self-Learning Loop

A scorer **improves its own rule**.

1. Capture outcome ("*market scorer ran hot*")
2. Extract a pattern from the bank
3. Write a rule to `lessons.md` , load next run

| *Session 4 (self-improvement half) — like `capture-lesson` .*

Demo one full loop: run → outcome → **new rule appears** → it changes the next score.

Your 32 Minutes

0-7 **Spec.** Brief + agent breakdown + Definition of Done, before building.

14-24 **Orchestrate.** Fan out to all criteria, weight, sum, map to verdict.

28-32 **Package.** Spec + agent defs + orchestration + one clean run for the 3-min pitch.

7-14 **Design scorers.** One scorer skill/agent + its description; get a valid `{score, reason}` .

24-28 **Verify.** DoD case + an edge idea (all 10s / all 1s / tie at 65). Fix it.

Judges hand you a fresh idea from a rotating set (Slack bot → Go · translation earbuds → De-risk First · a scopeless "AI for everything" pitch → Shelf). Boundary checks: all 10s → 100/100 → Go · all 1s → Shelf · a tie at exactly 65 confirms the ≥ 65 threshold.

Spec → Decompose → Orchestrate → Verify → Package.

Go Build + Challenge + Q&A



BREAKOUT ROOMS NOW

Don't type the scores — **orchestrate the agents that reason them out**. You're judged on the spec, the decomposition, and the proof you steered.



THIS WEEK'S MISSION

Apply the **orchestrator + specialists** pattern from today to one real task at work. Decompose it into sub-agents or skills the same way you decomposed the scoring criteria.

Bring the brief and the outcome to **Session 8 — Adoption Playbook & Retrospective**, the final session.



TO THE ROOM

- What was the **hardest part** of the orchestration design?